



Using nDPI over DPDK to Classify and Block Unwanted Network Traffic


Luca Deri <deri@ntop.org>
@lucaderi

Traffic Classification: an Overview



- Traffic classification is compulsory to understand the traffic flowing on a network and enhance user experience by tuning specific network parameters.
- Main classification methods include:
 - TCP/UDP port classification.
 - QoS based classification (DSCP).
 - Statistical Classification.
 - Deep Packet Inspection.

- Port-based Classification
 - In the early day of the Internet, network traffic protocols were identified by protocol and port.
 - Can classify only application protocols operating on well known ports (no rpcbind or portmap).
 - Easy to cheat and thus unreliable (TCP/80 != HTTP).
- QoS Markers (DSCP)
 - Similar to port classification but based on QoS tags.
 - Usually ignored as it is easy to cheat and forge.



- Classification of IP packets (size, port, flags, IP addresses) and flows (duration, frequency...).
- Based on rules written manually, or automatically using machine learning (ML) algorithms.
- ML requires a training set of very good quality, and it is generally computationally intensive.
- Detection rate can be as good as 95% for cases which were covered by the training set, and poor accuracy for all the other cases.

Deep Packet Inspection (DPI)



- Technique that inspects the packet payload.
- Computationally intensive with respect to simple packet header analysis.
- Concerns about privacy and confidentiality of inspected data.
- Encryption is becoming pervasive, thus challenging DPI techniques.
- No false positives unless statistical methods or IP range/flow analysis are used by DPI tools.

Using DPI in Traffic Monitoring



- Packet header analysis is no longer enough as it is unreliable and thus useless.
- Security and network administrators want to know what are the real protocols flowing on a network, this regardless of the port being used.
- Selective metadata extraction (e.g. HTTP URL or User-Agent) is necessary to perform accurate monitoring and thus this task should be performed by the DPI toolkit without replicating it on monitoring applications.

Welcome to nDPI



- In 2012 we decided to develop our own GNU LGPL DPI toolkit (based on a unmaintained project named OpenDPI) in order to build an open DPI layer for ntop and third-party applications (Wireshark, netfilter, ML tools...).
- Protocols supported exceed 240 and include:
 - P2P (Skype, BitTorrent)
 - Messaging (Viber, Whatsapp, MSN, Facebook)
 - Multimedia (YouTube, Last.gm, iTunes)
 - Conferencing (Webex, CitrixOnLine)
 - Streaming (Zattoo, Icecast, Shoutcast, Netflix)
 - Business (VNC, RDP, Citrix, *SQL)



What is a Protocol in nDPI? [1/2]



- Each protocol is identified as <major>.<minor> protocol. Example:
 - DNS.Facebook
 - QUIC.YouTube and **QUIC.YouTubeUpload**
- Caveat: Skype or Facebook are protocols in the nDPI world but not for IETF.
- The first question people ask when they have to evaluate a DPI toolkit is: how many protocol do you support? This is not the right question.

What is a Protocol in nDPI? [2/2]



- Today most protocols are HTTP/SSL-based.
- nDPI includes support for string-based protocols detection:
 - DNS query name
 - HTTP Host/Server header fields
 - SSL/QUIC SNI (Server Name Indication)
- Example: NetFlix detection


```
{ "netflix.com", NULL, "netflix" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },  
{ "nflxext.com", NULL, "nflxext" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },  
{ "nflximg.com", NULL, "nflximg" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },  
{ "nflximg.net", NULL, "nflximg" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },  
{ "nflxvideo.net", NULL, "nflxvideo" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },  
{ "nflxso.net", NULL, "nflxso" TLD, "NetFlix", NDPI_PROTOCOL_NETFLIX, NDPI_PROTOCOL_CATEGORY_STREAMING, NDPI_PROTOCOL_FUN },
```

nDPI Categories [1/2]



- Protocols are too many, and they increase daily.
- Many people are not familiar with protocol names.
- Often people ask us questions like “How can I prevent my children from using social networks?”
- Solution
 - nDPI allows protocols to be clustered in user-defined categories such as VoIP, P2P, Cloud...
 - Categories can include thousand of entries and can be (re-)loaded dynamically. Example: malware, mining, advertisement, banned site, inappropriate content...

nDPI Categories [2/2]




- Applications using nDPI are responsible for
 - Capturing (forwarding in inline mode) packets
 - Maintaining flow state.
- Based on flow protocol/port all dissectors that can potentially match the flow are applied sequentially starting from the one that most likely match.
- Each dissector is coded into a different .c file for the sake of modularity and extensibility.
- There is an extra .c file for IP matching (e.g. identify Spotify traffic based on Spotify AS).

Traffic Classification Lifecycle




- Based on traffic type (e.g. UDP traffic) dissectors are applied sequentially starting with the one that will most likely match the flow (e.g. for TCP/80 the HTTP dissector is tried first).
- Each flow maintains the state for non-matching dissectors in order to skip them in future iterations.
- Analysis lasts until a match is found or after too many attempts (8 packets is the upper-bound in our experience).

nDPI-based Applications: Architecture



Flow Lifecycle [1/2]



- DPI-oriented applications have to deal with flows
- A flow is identified by 5+1 tuple (VLAN, proto, IP/port src/dst).
- It is first created when the first packet is received
- Expires based on timeout or termination (FIN/RST)
- Flow packets are nDPI-processed until the protocol is detected until a max number of iterations (unknown protocol).

Flow Lifecycle [2/2]



- Flows are usually kept in a hash table hashed with the 5-tuple.
- Nasty traffic (e.g. DNS) could cause several collisions that might drive overall the performance down.
- Performance is affected by both Mpps (DPDK) and number of concurrent flows.
- Adding DPI in existing applications (e.g. a traffic monitoring application) must pay attention to flow lifecycle as much as packet processing.

DPDK Integration [1/2]



- nDPI is packet-capture neutral (DPDK, PF_RING, netmap, pcap...)
- Inside nDPI/example there is an application named *ndpiReader* that demonstrates how to use the nDPI API when reading from pcap files and DPDK.

```
$ cd nDPI/example  
$ make -f Makefile.dpdk  
$ sudo ./build/ndpiReader -c 1 --vdev=net_pcap0,iface=en01 -- -v 1
```

DPDK Integration [2/2]



```
while(dpdk_run_capture) {
    struct rte_mbuf *bufs[BURST_SIZE];
    u_int16_t num = rte_eth_rx_burst(dpdk_port_id, 0, bufs, BURST_SIZE);
    u_int i;

    if(num == 0) {
        usleep(1);
        continue;
    }

    for(i = 0; i < PREFETCH_OFFSET && i < num; i++)
        rte_prefetch0(rte_pktmbuf_mtod(bufs[i], void *));

    for(i = 0; i < num; i++) {
        char *data = rte_pktmbuf_mtod(bufs[i], char *);
        int len = rte_pktmbuf_pkt_len(bufs[i]);
        struct pcap_pkthdr h;

        h.len = h.caplen = len;
        gettimeofday(&h.ts, NULL);

        ndpi_process_packet((u_char*)&thread_id, &h, (const u_char *)data);
        rte_pktmbuf_free(bufs[i]);
    }
}
```

nDPI-over-DPDK Inline Mode



- You can take any DPDK application and add nDPI support to it

Branch: master ▾

[dpdk / examples / skeleton / basicfwd.c](#)

```
for (;;) {
    RTE_ETH_FOREACH_DEV(port) {

        /* Get burst of RX packets, from first port of pair. */
        struct rte_mbuf *bufs[BURST_SIZE];
        const uint16_t nb_rx = rte_eth_rx_burst(port, 0, bufs, BURST_SIZE);


        if (unlikely(nb_rx == 0))
            continue;

* nDPI processing code goes here *

        /* Send burst of TX packets, to second port of pair. */
        const uint16_t nb_tx = rte_eth_tx_burst(port ^ 1, 0, bufs, nb_rx);

        /* Free any unsent packets. */
        if (unlikely(nb_tx < nb_rx)) {
            uint16_t buf;
            for (buf = nb_tx; buf < nb_rx; buf++)
                rte_pktmbuf_free(bufs[buf]);
        }
    }
}
```

- PF_RING FT is natively integrated with nDPI for providing L7 protocol information
- The application does not need to deal directly with the nDPI library, as it:
 1. enables L7 detection through the API
 2. reads the L7 protocol from the exported metadata




nDPI + PF_RING FT + DPDK [2/3]



```
pfring_ft_table *ft = pfring_ft_create_table(  
    flags, max_flows, flow_idle_timeout, flow_lifetime_timeout);  
  
/* Callback for 'new flow' events */  
pfring_ft_set_new_flow_callback(ft, new_flow_callback, user);  
  
/* Callback for 'packet processed/classified' events */  
pfring_ft_set_flow_packet_callback(ft, packet_processed_callback, user);  
  
/* Callback for 'flow to be exported' events */  
pfring_ft_set_flow_export_callback(ft, export_flow_callback, user);  
  
...  
  
/* Process Captured Packets */  
while (1) {  
    int num = rte_eth_rx_burst(port_id, 0, bufs, BURST_SIZE);  
    pfring ft pcap_pkthdr h;  
    pfring ft ext_pkthdr ext_hdr = { 0 };  
  
    for (i = 0; i < num; i++) {  
        char *data = rte_pktmbuf_mtod(bufs[i], char *);  
        int len = rte_pktmbuf_pkt_len(bufs[i]);  
  
        if(pfring_ft_process(ft, (const u_char *)data, &h, &ext_hdr) != PFRING_FT_ACTION_DISCARD)  
            rte_eth_tx_burst(twin_port_id, 0, &bufs[i], 1);  
    }  
}
```

Full Example: https://github.com/ntop/PF_RING/blob/dev/userland/examples_ft/ftflow_dpdk.c

nDPI + PF_RING FT + DPDK [3/3]



nDPI: Packet Processing Performance: Pcap



nDPI Memory statistics:

nDPI Memory (once):	203.62 KB
Flow Memory (per flow):	2.01 KB
Actual Memory:	95.60 MB
Peak Memory:	95.60 MB
Setup Time:	1001 msec
Packet Processing Time:	813 msec

Traffic statistics:

Ethernet bytes:	1090890957	(includes ethernet CRC/IFC/trailer)
Discarded bytes:	247801	
IP packets:	1482145	of 1483237 packets total
IP bytes:	1055319477	(avg pkt size 711 bytes)
Unique flows:	36703	
TCP Packets:	1338624	
UDP Packets:	143521	
VLAN Packets:	0	
MPLS Packets:	0	
PPPoE Packets:	0	
Fragmented Packets:	1092	
Max Packet size:	1480	
Packet Len < 64:	590730	
Packet Len 64-128:	67824	
Packet Len 128-256:	66380	
Packet Len 256-1024:	157623	
Packet Len 1024-1500:	599588	
Packet Len > 1500:	0	
nDPI throughput:	1.82 M pps / 9.99 Gb/sec	 Single Core (E3 1241v3)
Analysis begin:	04/Aug/2010 04:15:23	
Analysis end:	04/Aug/2010 18:31:30	
Traffic throughput:	28.85 pps / 165.91 Kb/sec	
Traffic duration:	51367.223 sec	
Guessed flow protos:	0	

- 10 Gbit tests on Intel E3-1230 v5 3.4GHz DDR4 2133
- 100 Gbit tests on 2x Intel E5-2630 v2 2.6GHz DDR3 1600 (much slower than modern Xeon Scalable)
- nDPI integrated in a flow monitoring application (nProbe Cento)

Traffic	Capture Card	Number of Cores	Per Core Performance	All Cores Performance
10 Gbit / 64-byte packets	Intel 10G (X520)	1	14.8 Mpps / 10 Gbps	14.8 Mpps / 10 Gbps
100 Gbit / 1-kbyte packets	FPGA 100G	1	10.8 Mpps / 90 Gbps	10.8 Mpps / 90 Gbps
100 Gbit / 1-kbyte packets	FPGA 100G	4	2.8 Mpps / 24 Gbps	11.5 Mpps / 96 Gbps
100 Gbit / 64-byte packets	FPGA 100G	4	11.2 Mpps / 7.6 Gbps	45.2 Mpps / 30.4 Gbps
100 Gbit / 64-byte packets	FPGA 100G	6 + 6 (2 CPUs)	10.8 Mpps / 7.3 Gbps	130 Mpps / 87.6 Gbps

- Hyperscan is a high-performance regex matching library that can be used in nDPI instead of the native Aho-Corasick (configure --with-hyperscan)
- String matching is used in protocol detection.

HyperScan

nDPI Memory statistics:

nDPI Memory (once) :	203.62 KB
Flow Memory (per flow) :	2.01 KB
Actual Memory:	95.60 MB
Peak Memory:	95.60 MB
Setup Time:	1001 msec
Packet Processing Time:	813 msec

Aho-Corasick

nDPI Memory statistics:

nDPI Memory (once) :	203.62 KB
Flow Memory (per flow) :	2.01 KB
Actual Memory:	95.61 MB
Peak Memory:	95.61 MB
Setup Time:	11 msec
Packet Processing Time:	835 msec

Note: same test of slide 23 with HyperScan and Aho-Corasick

Evaluating nDPI



- nDPI has been evaluated both in terms of accuracy and performance.
- “The best accuracy we obtained from nDPI (91 points), PACE (82 points), UPC MLA (79 points), and Libprotoident (78 points)”
- Source: T. Bujlow, V. Carela-Español, P. Barlet-Ros, Comparison of Deep Packet Inspection (DPI) Tools for Traffic Classification, Technical Report, June 2013.

Final Remarks



- We have presented nDPI an open source DPI toolkit able to detect many popular Internet protocols and scale at 10 Gbit on commodity hardware platforms.
- Its open design make it suitable for using it both in open-source and security applications where code inspection is compulsory.
- Code Availability (GNU LGPLv3)
<https://github.com/ntop/nDPI>

Acknowledgment



- I would like to thank the Intel Software Innovator Program for supporting the development of nDPI

